

```
from google.colab import drive
drive.mount('/content/drive')

import os

BASE_DIR = "/content/drive/MyDrive/SkinAI_Project"
PROCESSED_DIR = os.path.join(BASE_DIR, "data", "processed")
MODELS_DIR = os.path.join(BASE_DIR, "models")

REPORTS_DIR = os.path.join(BASE_DIR, "reports")
FIG_DIR = os.path.join(REPORTS_DIR, "figures")
METRICS_DIR = os.path.join(REPORTS_DIR, "metrics")

os.makedirs(MODELS_DIR, exist_ok=True)
os.makedirs(FIG_DIR, exist_ok=True)
os.makedirs(METRICS_DIR, exist_ok=True)

print("BASE_DIR :", BASE_DIR)
print("PROCESSED_DIR:", PROCESSED_DIR)
print("MODELS_DIR :", MODELS_DIR)
print("Processed files:", os.listdir(PROCESSED_DIR))
```

```
Mounted at /content/drive
BASE_DIR : /content/drive/MyDrive/SkinAI_Project
PROCESSED_DIR: /content/drive/MyDrive/SkinAI_Project/data/processed
MODELS_DIR : /content/drive/MyDrive/SkinAI_Project/models
Processed files: ['text_augmented.csv', 'test_predictions.csv', 'structured', 'disease', 'unlabeled_text.csv', 'text_merged.csv']
```

```
import pandas as pd

train_path = os.path.join(PROPESSED_DIR, "train.csv")
val_path   = os.path.join(PROPESSED_DIR, "val.csv")
test_path  = os.path.join(PROPESSED_DIR, "test.csv")

df_train = pd.read_csv(train_path)
df_val   = pd.read_csv(val_path)
df_test  = pd.read_csv(test_path)

print("Train shape:", df_train.shape)
print("Val shape : ", df_val.shape)
print("Test shape :", df_test.shape)
print("Columns:", df_train.columns.tolist())

display(df_train.head())
print("\nLabel distribution (train top 30):")
print(df_train["label"].value_counts().head(30))

X_train = df_train["text"].astype(str).tolist()
y_train = df_train["label"].astype(str).tolist()

X_val = df_val["text"].astype(str).tolist()
y_val = df_val["label"].astype(str).tolist()

X_test = df_test["text"].astype(str).tolist()
y_test = df_test["label"].astype(str).tolist()

print("\nClasses:", df_train["label"].nunique())
```

```
Train shape: (7433, 3)
Val shape : (1593, 3)
Test shape : (1593, 3)
Columns: ['text', 'label', 'source']
```

		text	label	source
0	age: 68 gender: male symptom text: lately,...	Contact dermatitis (Irritant)	synthetic/clinical_cases_10000.csv	
1	age: 56 gender: male symptom text: recentl...	Tinea cruris (Jock itch)	synthetic/clinical_cases_10000.csv	
2	age: 48 gender: male symptom text: these d...	Cellulitis	synthetic/clinical_cases_10000.csv	
3	age: 35 gender: male symptom text: i've no...	Seborrheic dermatitis	synthetic/clinical_cases_10000.csv	
4	age: 48 gender: female symptom text: recen...	Impetigo	synthetic/clinical_cases_10000.csv	

Label distribution (train top 30):

label	count
Tinea corporis (Ringworm)	275
Urticaria (Hives)	269
Impetigo	268
Seborrheic dermatitis	264
Dyshidrotic eczema	257
Herpes simplex (Cold sores)	253
Rosacea	253
Pityriasis versicolor	252
Lichen planus	248
Folliculitis	248
Acne vulgaris	248
Tinea capitis (Scalp ringworm)	246
Contact dermatitis (Irritant)	246
Perioral dermatitis	246
Tinea cruris (Jock itch)	245
Molluscum contagiosum	245
Chronic urticaria	244
Vitiligo	243
Nummular eczema	242
Scabies	239
Acne rosacea	236
Cellulitis	234
Psoriasis vulgaris	234
Contact dermatitis (Allergic)	231
Herpes zoster (Shingles)	230
Atopic dermatitis (Eczema)	228
Photodermatitis	226
Warts (Verruca vulgaris)	223
Tinea pedis (Athlete's foot)	207
Dermatofibroma	28

Name: count, dtype: int64

Classes: 47

```
from sklearn.feature_extraction.text import TfidfVectorizer

vectorizer = TfidfVectorizer(
    ngram_range=(1, 2),
    max_features=5000,
    min_df=2,
    max_df=0.9
)

X_train_tfidf = vectorizer.fit_transform(X_train)
X_val_tfidf = vectorizer.transform(X_val)
X_test_tfidf = vectorizer.transform(X_test)

print("TF-IDF shapes:")
print(" Train:", X_train_tfidf.shape)
print(" Val : ", X_val_tfidf.shape)
print(" Test :", X_test_tfidf.shape)
```

```
TF-IDF shapes:
Train: (7433, 4572)
Val : (1593, 4572)
Test : (1593, 4572)
```

```

import numpy as np
from sklearn.linear_model import SGDClassifier
from sklearn.metrics import log_loss
from sklearn.utils.class_weight import compute_class_weight

# Classes must be fixed across partial_fit calls
classes = np.unique(y_train)

#  SGD model WITHOUT class_weight="balanced" (not supported in partial_fit)
sgd = SGDClassifier(
    loss="log_loss",
    max_iter=1,           # not used in partial_fit loop, but keep valid
    tol=None,             # important: disable early stopping for partial_fit epochs
    random_state=42,
    learning_rate="optimal"
)

#  Compute class weights manually (balanced) and convert to sample weights
class_weights = compute_class_weight(
    class_weight="balanced",
    classes=classes,
    y=np.array(y_train)
)
class_weight_dict = {c: w for c, w in zip(classes, class_weights)}

# sample_weight for each training example
sample_weight_train = np.array([class_weight_dict[label] for label in y_train], dtype=float)

print("Computed class weights (balanced):")
for c in classes[:10]:
    print(f" {c}: {class_weight_dict[c]:.4f}")
print("... (showing first 10)")

train_losses = []
val_losses = []

EPOCHS = 25

for epoch in range(1, EPOCHS + 1):
    # partial_fit supports sample_weight
    if epoch == 1:
        sgd.partial_fit(X_train_tfidf, y_train, classes=classes, sample_weight=sample_weight_train)
    else:
        sgd.partial_fit(X_train_tfidf, y_train, sample_weight=sample_weight_train)

    # Compute log loss using probabilities
    train_probs = sgd.predict_proba(X_train_tfidf)
    val_probs   = sgd.predict_proba(X_val_tfidf)

    tr_loss = log_loss(y_train, train_probs, labels=classes)
    va_loss = log_loss(y_val, val_probs, labels=classes)

    train_losses.append(tr_loss)
    val_losses.append(va_loss)

    print(f"Epoch {epoch:02d} | Train LogLoss: {tr_loss:.4f} | Val LogLoss: {va_loss:.4f}")

print("\nTraining finished (SGD partial_fit + balanced sample weights).")

```

```

Computed class weights (balanced):
Acne: 7.5309
Acne rosacea: 0.6701
Acne vulgaris: 0.6377
Actinic keratosis: 6.3260
Athlete Foot (Tinea Pedis): 19.7686
Athlete's Foot (Tinea Pedis): 19.7686
Atopic Dermatitis: 6.3260
Atopic dermatitis (Eczema): 0.6936
Benign keratosis: 5.6482
Cellulitis: 0.6759
... (showing first 10)
Epoch 01 | Train LogLoss: 0.8857 | Val LogLoss: 0.9934
Epoch 02 | Train LogLoss: 0.8564 | Val LogLoss: 0.9718
Epoch 03 | Train LogLoss: 0.8498 | Val LogLoss: 0.9660
Epoch 04 | Train LogLoss: 0.8462 | Val LogLoss: 0.9625
Epoch 05 | Train LogLoss: 0.8438 | Val LogLoss: 0.9601
Epoch 06 | Train LogLoss: 0.8418 | Val LogLoss: 0.9582

```

```

Epoch 07 | Train LogLoss: 0.8403 | Val LogLoss: 0.9566
Epoch 08 | Train LogLoss: 0.8389 | Val LogLoss: 0.9553
Epoch 09 | Train LogLoss: 0.8378 | Val LogLoss: 0.9542
Epoch 10 | Train LogLoss: 0.8368 | Val LogLoss: 0.9533
Epoch 11 | Train LogLoss: 0.8359 | Val LogLoss: 0.9524
Epoch 12 | Train LogLoss: 0.8352 | Val LogLoss: 0.9517
Epoch 13 | Train LogLoss: 0.8344 | Val LogLoss: 0.9510
Epoch 14 | Train LogLoss: 0.8338 | Val LogLoss: 0.9503
Epoch 15 | Train LogLoss: 0.8332 | Val LogLoss: 0.9497
Epoch 16 | Train LogLoss: 0.8326 | Val LogLoss: 0.9492
Epoch 17 | Train LogLoss: 0.8321 | Val LogLoss: 0.9487
Epoch 18 | Train LogLoss: 0.8316 | Val LogLoss: 0.9482
Epoch 19 | Train LogLoss: 0.8312 | Val LogLoss: 0.9478
Epoch 20 | Train LogLoss: 0.8307 | Val LogLoss: 0.9474
Epoch 21 | Train LogLoss: 0.8303 | Val LogLoss: 0.9470
Epoch 22 | Train LogLoss: 0.8299 | Val LogLoss: 0.9466
Epoch 23 | Train LogLoss: 0.8296 | Val LogLoss: 0.9462
Epoch 24 | Train LogLoss: 0.8292 | Val LogLoss: 0.9459
Epoch 25 | Train LogLoss: 0.8289 | Val LogLoss: 0.9456

```

Training finished (SGD partial_fit + balanced sample weights).


```

import matplotlib.pyplot as plt

plt.figure(figsize=(8,5))
plt.plot(range(1, len(train_losses)+1), train_losses, marker="o", label="Train LogLoss")
plt.plot(range(1, len(val_losses)+1), val_losses, marker="o", label="Val LogLoss")
plt.title("SGD (log_loss) - Training Loss Curve")
plt.xlabel("Epoch")
plt.ylabel("Log Loss (Lower is better)")
plt.grid(True, alpha=0.3)
plt.legend()

out_path = os.path.join(FILE_DIR, "sgd_logloss_loss_curve.png")
plt.savefig(out_path, dpi=200, bbox_inches="tight")
plt.show()
print("Saved:", out_path)

```



Saved: /content/drive/MyDrive/SkinAI_Project/reports/figures/sgd_logloss_loss_curve.png

```

from sklearn.metrics import accuracy_score, f1_score, classification_report

def evaluate_split(name, model, X, y_true):
    y_pred = model.predict(X)
    acc = accuracy_score(y_true, y_pred)
    f1m = f1_score(y_true, y_pred, average="macro")
    print(f"\n{n} == {name} ==")
    print("Accuracy:", round(acc, 4))
    print("Macro-F1:", round(f1m, 4))
    print("\nClassification report:")
    print(classification_report(y_true, y_pred, zero_division=0))
    return y_pred, acc, f1m

```

```
y_pred_train, acc_train, f1_train = evaluate_split("TRAIN", sgd, X_train_tfidf, y_train)
y_pred_val, acc_val, f1_val = evaluate_split("VALIDATION", sgd, X_val_tfidf, y_val)
y_pred_test, acc_test, f1_test = evaluate_split("TEST", sgd, X_test_tfidf, y_test)

report_path = os.path.join(METRICS_DIR, "sgd_logloss_report.txt")
with open(report_path, "w", encoding="utf-8") as f:
    f.write(classification_report(y_test, y_pred_test, zero_division=0))
print("Saved report:", report_path)
```


==== TRAIN ====
Accuracy: 0.7999
Macro-F1: 0.853

Classification report:

	precision	recall	f1-score	support
Acne	0.78	1.00	0.88	21
Acne rosacea	1.00	1.00	1.00	236
Acne vulgaris	1.00	0.99	1.00	248
Actinic keratosis	1.00	1.00	1.00	25
Athlete Foot (Tinea Pedis)	1.00	1.00	1.00	8
Athlete's Foot (Tinea Pedis)	1.00	1.00	1.00	8
Atopic Dermatitis	0.93	1.00	0.96	25
Atopic dermatitis (Eczema)	0.30	0.33	0.32	228
Benign keratosis	0.88	1.00	0.93	28
Cellulitis	1.00	1.00	1.00	234
Chronic urticaria	1.00	1.00	1.00	244
Contact Dermatitis	1.00	1.00	1.00	10
Contact dermatitis (Allergic)	0.26	0.32	0.29	231
Contact dermatitis (Irritant)	1.00	1.00	1.00	246
Dermatofibroma	1.00	1.00	1.00	28
Dyshidrotic eczema	1.00	1.00	1.00	257
Eczema	0.75	1.00	0.86	15
Folliculitis	0.37	0.32	0.34	248
Herpes simplex (Cold sores)	1.00	1.00	1.00	253
Herpes zoster (Shingles)	0.28	0.32	0.30	230
Hives (Urticaria)	1.00	1.00	1.00	15
Impetigo	1.00	1.00	1.00	268
Lichen planus	1.00	1.00	1.00	248
Melanocytic nevus	0.96	1.00	0.98	25
Melanoma	1.00	1.00	1.00	25
Molluscum contagiosum	1.00	1.00	1.00	245
Nummular eczema	1.00	1.00	1.00	242
Perioral dermatitis	1.00	1.00	1.00	246
Photodermatitis	0.26	0.22	0.24	226
Pityriasis versicolor	1.00	1.00	1.00	252
Psoriasis	0.94	1.00	0.97	15
Psoriasis vulgaris	1.00	1.00	1.00	234
Ringworm (Tinea Corporis)	1.00	1.00	1.00	15
Rosacea	1.00	0.97	0.99	253
Scabies	0.34	0.29	0.31	239
Seborrheic dermatitis	1.00	1.00	1.00	264
Shingles (Herpes Zoster)	1.00	1.00	1.00	15
Squamous cell carcinoma	1.00	1.00	1.00	25
Tinea Ringworm Candidiasis	0.96	1.00	0.98	25
Tinea capitis (Scalp ringworm)	1.00	1.00	1.00	246
Tinea corporis (Ringworm)	1.00	1.00	1.00	275
Tinea cruris (Jock itch)	1.00	1.00	1.00	245
Tinea pedis (Athlete's foot)	0.21	0.28	0.24	207
Urticaria (Hives)	1.00	1.00	1.00	269
Vascular lesion	0.93	1.00	0.96	25
Vitiligo	0.32	0.26	0.29	243
Warts (Verruca vulgaris)	0.29	0.25	0.27	223
accuracy			0.80	7433
macro avg	0.85	0.86	0.85	7433

```
plt.figure(figsize=(7,5))
plt.bar(["Train", "Validation", "Test"], [acc_train, acc_val, acc_test])
plt.title("Accuracy Comparison (SGD log_loss)")
plt.ylabel("Accuracy")
plt.ylim(0, 1)
plt.grid(True, axis="y", alpha=0.3)

out_path = os.path.join(FIG_DIR, "sgd_logloss_accuracy_bar.png")
plt.savefig(out_path, dpi=200, bbox_inches="tight")
plt.show()
print("Saved:", out_path)
```




Saved: /content/drive/MyDrive/SkinAI_Project/reports/figures/sgd_logloss_accuracy_bar.png

```
from sklearn.metrics import f1_score
import pandas as pd

unique_labels = df_test["label"].unique()
f1_scores = {label: f1_score(y_test, y_pred_test, labels=[label], average="macro") for label in unique_labels}

pd.Series(f1_scores).sort_values(ascending=False).head(15).plot(kind="bar", figsize=(10,5))
plt.title("Top 15 F1-scores by Disease (SGD log_loss)")
plt.ylabel("F1-score")
plt.xticks(rotation=45, ha="right")
plt.grid(True, axis="y", alpha=0.3)

out_path = os.path.join(FIG_DIR, "sgd_logloss_top15_f1.png")
plt.savefig(out_path, dpi=200, bbox_inches="tight")
plt.show()
print("Saved:", out_path)
```



```
import numpy as np
from sklearn.metrics import confusion_matrix

N = 10
top_labels = pd.Series(y_test).value_counts().head(N).index.tolist()
```

```

mask = [y in top_labels for y in y_test]
y_true_top = np.array(y_test)[mask]
y_pred_top = np.array(y_pred_test)[mask]

cm = confusion_matrix(y_true_top, y_pred_top, labels=top_labels)

fig, ax = plt.subplots(figsize=(10, 8))
im = ax.imshow(cm, interpolation="nearest")
ax.figure.colorbar(im, ax=ax)


ax.set(
    xticks=np.arange(len(top_labels)),
    yticks=np.arange(len(top_labels)),
    xticklabels=top_labels,
    yticklabels=top_labels,
    ylabel="True label",
    xlabel="Predicted label",
    title=f"Confusion Matrix (Top {N} labels, Test set) - SGD log_loss"
)
plt.setp(ax.get_xticklabels(), rotation=45, ha="right")

thresh = cm.max() / 2. if cm.max() > 0 else 0
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        val = cm[i, j]
        ax.text(j, i, str(val), ha="center", va="center",
                color="white" if val > thresh else "black", fontsize=9)

plt.tight_layout()

out_path = os.path.join(FIG_DIR, "sgd_logloss_confusion_top10.png")
plt.savefig(out_path, dpi=200, bbox_inches="tight")
plt.show()
print("Saved:", out_path)

```



```

from sklearn.metrics import log_loss
from sklearn.preprocessing import label_binarize

test_probs = sgd.predict_proba(X_test_tfidf)

test_logloss = log_loss(y_test, test_probs, labels=classes)
print("TEST Log Loss:", round(test_logloss, 4))

```

```


y_test_bin = label_binarize(y_test, classes=classes)
brier = np.mean(np.sum((y_test_bin - test_probs)**2, axis=1))
print("TEST Brier Score:", round(float(brier), 4))

# Plot
plt.figure(figsize=(7,4))
plt.bar(["LogLoss", "Brier"], [test_logloss, brier])
plt.title("Probability Quality Metrics (SGD log_loss) - Test")
plt.grid(True, axis="y", alpha=0.3)

out_path = os.path.join(FIG_DIR, "sgd_logloss_prob_quality.png")
plt.savefig(out_path, dpi=200, bbox_inches="tight")
plt.show()
print("Saved:", out_path)

```

TEST Log Loss: 0.9595
TEST Brier Score: 0.3287



Saved: /content/drive/MyDrive/SkinAI_Project/reports/figures/sgd_logloss_prob_quality.png

```

from sklearn.calibration import calibration_curve


conf_test = test_probs.max(axis=1)
correct_test = (np.array(y_pred_test) == np.array(y_test)).astype(int)

frac_pos, mean_pred = calibration_curve(correct_test, conf_test, n_bins=10, strategy="uniform")

plt.figure(figsize=(6,5))
plt.plot(mean_pred, frac_pos, marker="o", label="Model")
plt.plot([0, 1], [0, 1], linestyle="--", label="Perfect")
plt.title("Calibration Curve (Test) - SGD log_loss")
plt.xlabel("Mean predicted confidence")
plt.ylabel("Fraction correct")
plt.grid(True, alpha=0.3)
plt.legend()

out_path = os.path.join(FIG_DIR, "sgd_logloss_calibration_curve.png")
plt.savefig(out_path, dpi=200, bbox_inches="tight")
plt.show()
print("Saved:", out_path)


```



Saved: /content/drive/MyDrive/SkinAI_Project/reports/figures/sgd_logloss_calibration_curve.png

```
plt.figure(figsize=(7,4))
plt.hist(conf_test * 100, bins=20)
plt.title("Confidence Distribution (Test) - SGD log_loss")
plt.xlabel("Confidence %")
plt.ylabel("Count")
plt.grid(True, alpha=0.3)

out_path = os.path.join(FIG_DIR, "sgd_logloss_confidence_hist.png")
plt.savefig(out_path, dpi=200, bbox_inches="tight")
plt.show()
print("Saved:", out_path)
```



Saved: /content/drive/MyDrive/SkinAI_Project/reports/figures/sgd_logloss_confidence_hist.png

```

import joblib
Saved classifier to: /content/drive/MyDrive/SkinAI_Project/models/skin_text_classifier_tfidf_sgd_logloss.joblib
Saved vectorizer to: /content/drive/MyDrive/SkinAI_Project/models/vectorizer_tfidf.joblib
MODEL_PATH = os.path.join(MODELS_DIR, "skin_text_classifier_tfidf_sgd_logloss.joblib")
VECT_PATH = os.path.join(MODELS_DIR, "vectorizer_tfidf.joblib")

def predict_with_confidence(model, vectorizer, text: str):
    """
    Predict label + confidence for SGD (log_loss).
    confidence = predicted class probability (0..1) if predict_proba exists.
    """
    X = vectorizer.transform([text])
    pred = model.predict(X)[0]

    if hasattr(model, "predict_proba"):
        proba = model.predict_proba(X)[0]
        classes = list(model.classes_)
        conf = float(proba[classes.index(pred)])
        return pred, conf

    return pred, None

samples = [
    "I have pimples and red bumps with white pus on my cheeks and forehead.",
    "My skin is very dry and itchy with red patches that crack sometimes.",
    "I have a circular itchy rash with raised red border and clear center."
]

for s in samples:
    p, c = predict_with_confidence(sgd, vectorizer, s) # ✅ use sgd here
    print("\nTEXT:", s)
    print("PRED:", p)
    print("CONF:", "N/A" if c is None else f"{c*100:.2f}%")

```

```

TEXT: I have pimples and red bumps with white pus on my cheeks and forehead.
PRED: Acne vulgaris
CONF: 10.05%

TEXT: My skin is very dry and itchy with red patches that crack sometimes.
PRED: Benign keratosis
CONF: 4.72%

TEXT: I have a circular itchy rash with raised red border and clear center.
PRED: Tinea corporis (Ringworm)
CONF: 11.21%

```

```

import numpy as np

def predict_topk(model, vectorizer, text: str, k: int = 3):
    X = vectorizer.transform([text])

    if not hasattr(model, "predict_proba"):
        pred = model.predict(X)[0]
        return [(pred, None)]

    proba = model.predict_proba(X)[0]
    classes = np.array(model.classes_)
    top_idx = np.argsort(proba)[-1:][:-k]

    return [(classes[i], float(proba[i])) for i in top_idx]

for s in samples:
    top3 = predict_topk(sgd, vectorizer, s, k=3)
    print("\nTEXT:", s)
    for rank, (lbl, conf) in enumerate(top3, 1):
        print(f"Top-{rank}: {lbl} | {conf*100:.2f}%")

```

```

TEXT: I have pimples and red bumps with white pus on my cheeks and forehead.
Top-1: Acne vulgaris | 10.05%
Top-2: Acne rosacea | 9.52%
Top-3: Folliculitis | 5.31%

```