

```

from google.colab import drive
drive.mount('/content/drive')

import os

# Base project directory on Google Drive
BASE_DIR = "/content/drive/MyDrive/SkinAI_Project"

# Data dirs
DATA_DIR = f"{BASE_DIR}/data"
RAW_DIR = f"{DATA_DIR}/raw"
PROCESSED_DIR = f"{DATA_DIR}/processed"

# Models + notebooks
MODELS_DIR = f"{BASE_DIR}/models"

# Outputs for graphs/reports (create if not exist)
REPORTS_DIR = f"{BASE_DIR}/reports"
FIG_DIR = f"{REPORTS_DIR}/figures"
METRICS_DIR = f"{REPORTS_DIR}/metrics"

os.makedirs(MODELS_DIR, exist_ok=True)
os.makedirs(FIG_DIR, exist_ok=True)
os.makedirs(METRICS_DIR, exist_ok=True)

print("BASE_DIR : ", BASE_DIR)
print("PROCESSED_DIR:", PROCESSED_DIR)
print("MODELS_DIR : ", MODELS_DIR)

print("\nProcessed dir contents:")
print(os.listdir(PROCESSED_DIR))

```

```

Mounted at /content/drive
BASE_DIR : /content/drive/MyDrive/SkinAI_Project
PROCESSED_DIR: /content/drive/MyDrive/SkinAI_Project/data/processed
MODELS_DIR : /content/drive/MyDrive/SkinAI_Project/models

Processed dir contents:
['text_augmented.csv', 'test_predictions.csv', 'structured', 'disease', 'unlabeled_text.csv', 'text_merged.csv', 'text_cleaned.csv']

```

```

import pandas as pd
import os

train_path = os.path.join(PROPESSED_DIR, "train.csv")
val_path   = os.path.join(PROPESSED_DIR, "val.csv")
test_path  = os.path.join(PROPESSED_DIR, "test.csv")

df_train = pd.read_csv(train_path)
df_val   = pd.read_csv(val_path)
df_test  = pd.read_csv(test_path)

print("Train shape:", df_train.shape)
print("Val shape : ", df_val.shape)
print("Test shape : ", df_test.shape)

print("\nTrain columns:", df_train.columns.tolist())
print("Val columns : ", df_val.columns.tolist())
print("Test columns : ", df_test.columns.tolist())

print("\nSample train rows:")
display(df_train.head())

print("Null counts (train):")
print(df_train.isna().sum())

print("\nLabel distribution (train, top 30):")
print(df_train["label"].value_counts().head(30))

num_classes = df_train["label"].nunique()
print("\nNumber of unique labels in train:", num_classes)

```

```

Train shape: (7433, 3)
Val shape : (1593, 3)
Test shape : (1593, 3)

Train columns: ['text', 'label', 'source']
Val columns : ['text', 'label', 'source']
Test columns : ['text', 'label', 'source']

Sample train rows:

```

	text	label	source
0	age: 68 gender: male symptom text: lately,...	Contact dermatitis (Irritant)	synthetic/clinical_cases_10000.csv
1	age: 56 gender: male symptom text: recentl...	Tinea cruris (Jock itch)	synthetic/clinical_cases_10000.csv
2	age: 48 gender: male symptom text: these d...	Cellulitis	synthetic/clinical_cases_10000.csv
3	age: 35 gender: male symptom text: i've no...	Seborrheic dermatitis	synthetic/clinical_cases_10000.csv
4	age: 48 gender: female symptom text: recen...	Impetigo	synthetic/clinical_cases_10000.csv

Null counts (train):

text	0
label	0
source	0
dtype:	int64

Label distribution (train, top 30):

label	count
Tinea corporis (Ringworm)	275
Urticaria (Hives)	269
Impetigo	268
Seborrheic dermatitis	264
Dyshidrotic eczema	257
Herpes simplex (Cold sores)	253
Rosacea	253
Pityriasis versicolor	252
Lichen planus	248
Folliculitis	248
Acne vulgaris	248
Tinea capitis (Scalp ringworm)	246
Contact dermatitis (Irritant)	246
Perioral dermatitis	246
Tinea cruris (Jock itch)	245
Molluscum contagiosum	245
Chronic urticaria	244
Vitiligo	243
Nummular eczema	242
Scabies	239
Acne rosacea	236
Cellulitis	234
Psoriasis vulgaris	234
Contact dermatitis (Allergic)	231
Herpes zoster (Shingles)	230
Atopic dermatitis (Eczema)	228
Photodermatitis	226
Warts (Verruca vulgaris)	223
Tinea pedis (Athlete's foot)	207
Dermatofibroma	28

```

# Features and labels
X_train = df_train["text"].astype(str).tolist()
y_train = df_train["label"].astype(str).tolist()

X_val = df_val["text"].astype(str).tolist()
y_val = df_val["label"].astype(str).tolist()

X_test = df_test["text"].astype(str).tolist()
y_test = df_test["label"].astype(str).tolist()

print("\nExamples:")
for i in range(3):
    print(f"[{i}] label={y_train[i]} | text={X_train[i][:120]}...")

```

Examples:

```

[0] label=Contact dermatitis (Irritant) | text=age: 68 | gender: male | symptom text: lately, for around 6 months, i've had redn
[1] label=Tinea cruris (Jock itch) | text=age: 56 | gender: male | symptom text: recently, for the last 5 weeks, i've had an itc
[2] label=Cellulitis | text=age: 48 | gender: male | symptom text: these days, for the last 7 weeks, i've had redness with swell

```

```
from sklearn.feature_extraction.text import TfidfVectorizer

vectorizer = TfidfVectorizer(
    ngram_range=(1, 2),    # unigrams + bigrams
    max_features=5000,     # tuneable
    min_df=2,              # ignore rare terms
    max_df=0.9             # ignore too frequent terms
)

# Fit on train only
X_train_tfidf = vectorizer.fit_transform(X_train)
X_val_tfidf   = vectorizer.transform(X_val)
X_test_tfidf  = vectorizer.transform(X_test)

print("TF-IDF shapes:")
print(" X_train:", X_train_tfidf.shape)
print(" X_val : ", X_val_tfidf.shape)
print(" X_test :", X_test_tfidf.shape)
```

```
TF-IDF shapes:
X_train: (7433, 4572)
X_val : (1593, 4572)
X_test : (1593, 4572)
```

```
from sklearn.svm import LinearSVC

# Base SVM model (no probabilities)
svm_base = LinearSVC(class_weight="balanced", random_state=42)

print("Training LinearSVC (base) on TRAIN...")
svm_base.fit(X_train_tfidf, y_train)
print("Training finished.")
```

```
Training LinearSVC (base) on TRAIN...
Training finished.
```

```
from sklearn.calibration import CalibratedClassifierCV

# Calibrate probabilities using validation set (prefit)
svm_cal = CalibratedClassifierCV(
    estimator=svm_base,
    method="sigmoid",
    cv="prefit"
)

print("Calibrating probabilities using VALIDATION set...")
svm_cal.fit(X_val_tfidf, y_val)
print("Calibration finished.")
```

```
Calibrating probabilities using VALIDATION set...
/usr/local/lib/python3.12/dist-packages/sklearn/calibration.py:333: UserWarning: The `cv='prefit'` option is deprecated in 1.6 a
  warnings.warn(
Calibration finished.
```

```
from sklearn.metrics import accuracy_score, classification_report

def evaluate_split(name, model, X, y_true):
    y_pred = model.predict(X)
    acc = accuracy_score(y_true, y_pred)
    print(f"\n{name} == {name} ==")
    print(f"Accuracy:{", round(acc, 4), ")")
    print("\nClassification report:")
    print(classification_report(y_true, y_pred, zero_division=0))
    return y_pred, acc

# Train
y_pred_train, acc_train = evaluate_split("TRAIN", svm_cal, X_train_tfidf, y_train)

# Validation
y_pred_val, acc_val = evaluate_split("VALIDATION", svm_cal, X_val_tfidf, y_val)

# Test
```

```
y_pred_test, acc_test = evaluate_split("TEST", svm_cal, X_test_tfidf, y_test)
```

==== TRAIN ====
Accuracy: 0.8087

Classification report:

		precision	recall	f1-score	support
	Acne	1.00	1.00	1.00	21
	Acne rosacea	1.00	1.00	1.00	236
	Acne vulgaris	1.00	1.00	1.00	248
	Actinic keratosis	1.00	1.00	1.00	25
	Athlete Foot (Tinea Pedis)	1.00	0.88	0.93	8
	Athlete's Foot (Tinea Pedis)	0.89	1.00	0.94	8
	Atopic Dermatitis	1.00	1.00	1.00	25
	Atopic dermatitis (Eczema)	0.34	0.32	0.33	228
	Benign keratosis	1.00	1.00	1.00	28
	Cellulitis	1.00	1.00	1.00	234
	Chronic urticaria	1.00	1.00	1.00	244
	Contact Dermatitis	1.00	1.00	1.00	10
	Contact dermatitis (Allergic)	0.30	0.27	0.28	231
	Contact dermatitis (Irritant)	1.00	1.00	1.00	246
	Dermatofibroma	1.00	1.00	1.00	28
	Dyshidrotic eczema	1.00	1.00	1.00	257
	Eczema	1.00	1.00	1.00	15
	Folliculitis	0.38	0.37	0.37	248
	Herpes simplex (Cold sores)	1.00	1.00	1.00	253
	Herpes zoster (Shingles)	0.30	0.28	0.29	230
	Hives (Urticaria)	1.00	1.00	1.00	15
	Impetigo	1.00	1.00	1.00	268
	Lichen planus	1.00	1.00	1.00	248
	Melanocytic nevus	1.00	1.00	1.00	25
	Melanoma	1.00	1.00	1.00	25
	Molluscum contagiosum	1.00	1.00	1.00	245
	Nummular eczema	1.00	1.00	1.00	242
	Perioral dermatitis	1.00	1.00	1.00	246
	Photodermatitis	0.29	0.36	0.32	226
	Pityriasis versicolor	1.00	1.00	1.00	252
	Psoriasis	1.00	1.00	1.00	15
	Psoriasis vulgaris	1.00	1.00	1.00	234
	Ringworm (Tinea Corporis)	1.00	1.00	1.00	15
	Rosacea	1.00	1.00	1.00	253
	Scabies	0.31	0.41	0.35	239
	Seborrheic dermatitis	1.00	1.00	1.00	264
	Shingles (Herpes Zoster)	1.00	1.00	1.00	15
	Squamous cell carcinoma	1.00	1.00	1.00	25
	Tinea Ringworm Candidiasis	1.00	1.00	1.00	25
	Tinea capititis (Scalp ringworm)	1.00	1.00	1.00	246
	Tinea corporis (Ringworm)	1.00	1.00	1.00	275
	Tinea cruris (Jock itch)	1.00	1.00	1.00	245
	Tinea pedis (Athlete's foot)	0.27	0.24	0.25	207
	Urticaria (Hives)	1.00	1.00	1.00	269
	Vascular lesion	1.00	1.00	1.00	25
	Vitiligo	0.32	0.36	0.34	243
	Warts (Verruca vulgaris)	0.32	0.22	0.26	223
	accuracy			0.81	7433
	macro avg	0.87	0.87	0.87	7433
	weighted avg	0.81	0.81	0.81	7433

```


import matplotlib.pyplot as plt

accuracies = {
    "Train": acc_train,
    "Validation": acc_val,
    "Test": acc_test
}

plt.figure(figsize=(7,5))
plt.bar(accuracies.keys(), accuracies.values())
plt.title("Accuracy Comparison (SVM Calibrated)")
plt.ylabel("Accuracy")
plt.ylim(0, 1)
plt.grid(True, axis="y", alpha=0.3)

out_path = os.path.join(FILE_DIR, "svm_calibrated_accuracy_bar.png")
plt.savefig(out_path, dpi=200, bbox_inches="tight")
plt.show()
print("Saved:", out_path)

```



Saved: /content/drive/MyDrive/SkinAI_Project/reports/figures/svm_calibrated_accuracy_bar.png

```

from sklearn.metrics import f1_score
import pandas as pd


unique_labels = df_test["label"].unique()
f1_scores = {}

for label in unique_labels:
    f1_scores[label] = f1_score(y_test, y_pred_test, labels=[label], average="macro")

pd.Series(f1_scores).sort_values(ascending=False).head(15).plot(kind="bar", figsize=(10,5))
plt.title("Top 15 F1-scores by Disease (SVM Calibrated)")
plt.ylabel("F1-score")
plt.xticks(rotation=45, ha="right")
plt.grid(True, axis="y", alpha=0.3)

out_path = os.path.join(FILE_DIR, "svm_calibrated_top15_f1.png")
plt.savefig(out_path, dpi=200, bbox_inches="tight")
plt.show()
print("Saved:", out_path)

```



```

import numpy as np
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt

# Get top N labels by frequency in test set
N = 10
label_counts = pd.Series(y_test).value_counts().head(N)
top_labels = label_counts.index.tolist()

# Filter to top N labels
mask = [y in top_labels for y in y_test]
y_true_top = np.array(y_test)[mask]
y_pred_top = np.array(y_pred_test)[mask]

cm = confusion_matrix(y_true_top, y_pred_top, labels=top_labels)

fig, ax = plt.subplots(figsize=(10, 8))
im = ax.imshow(cm, interpolation="nearest")
ax.figure.colorbar(im, ax=ax)


ax.set(
    xticks=np.arange(len(top_labels)),
    yticks=np.arange(len(top_labels)),
    xticklabels=top_labels,
    yticklabels=top_labels,
    ylabel="True label",
    xlabel="Predicted label",
    title=f"Confusion Matrix (Top {N} labels, Test set) - SVM Calibrated"
)
plt.setp(ax.get_xticklabels(), rotation=45, ha="right", rotation_mode="anchor")

# Annotate cells (show ALL numbers including zeros if you want)
thresh = cm.max() / 2. if cm.max() > 0 else 0
for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        val = cm[i, j]
        ax.text(j, i, format(val, "d"),
                ha="center", va="center",
                color="white" if val > thresh else "black",
                fontsize=9
        )

plt.tight_layout()

out_path = os.path.join(FILE_DIR, "svm_calibrated_confusion_top10.png")
plt.savefig(out_path, dpi=200, bbox_inches="tight")
plt.show()
print("Saved:", out_path)

```



```

import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import log_loss
from sklearn.preprocessing import label_binarize

# Probabilities from calibrated SVM
val_probs = svm_cal.predict_proba(X_val_tfidf)
test_probs = svm_cal.predict_proba(X_test_tfidf)

# Log loss
val_logloss = log_loss(y_val, val_probs, labels=svm_cal.classes_)
test_logloss = log_loss(y_test, test_probs, labels=svm_cal.classes_)

print("VAL Log Loss :", round(val_logloss, 4))
print("TEST Log Loss:", round(test_logloss, 4))


plt.figure(figsize=(6,4))
plt.plot(["VAL", "TEST"], [val_logloss, test_logloss], marker="o")
plt.title("SVM Calibrated - Log Loss (Lower is better)")
plt.ylabel("Log Loss")
plt.grid(True, alpha=0.3)
out_path = os.path.join(FILE_DIR, "svm_calibrated_logloss.png")
plt.savefig(out_path, dpi=200, bbox_inches="tight")
plt.show()
print("Saved:", out_path)

# Brier score (multi-class)
y_test_bin = label_binarize(y_test, classes=svm_cal.classes_)
brier = np.mean(np.sum((y_test_bin - test_probs)**2, axis=1))
print("TEST Brier Score (Lower is better):", round(float(brier), 4))


plt.figure(figsize=(6,4))
plt.bar(["TEST"], [brier])
plt.title("SVM Calibrated - Brier Score (Lower is better)")
plt.ylabel("Brier Score")
plt.ylim(0, max(brier * 1.3, 0.05))
plt.grid(True, axis="y", alpha=0.3)
out_path = os.path.join(FILE_DIR, "svm_calibrated_brier.png")
plt.savefig(out_path, dpi=200, bbox_inches="tight")
plt.show()
print("Saved:", out_path)

```

VAL Log Loss : 0.7934
TEST Log Loss: 0.8255



Saved: /content/drive/MyDrive/SkinAI_Project/reports/figures/svm_calibrated_logloss.png
TEST Brier Score (Lower is better): 0.2902



Saved: /content/drive/MyDrive/SkinAI_Project/reports/figures/svm_calibrated_brier.png


```
from sklearn.calibration import calibration_curve
import numpy as np
import matplotlib.pyplot as plt

# Confidence = max predicted probability
conf_test = test_probs.max(axis=1)
correct_test = (np.array(y_pred_test) == np.array(y_test)).astype(int)

frac_pos, mean_pred = calibration_curve(correct_test, conf_test, n_bins=10, strategy="uniform")

plt.figure(figsize=(6,5))
plt.plot(mean_pred, frac_pos, marker="o", label="Model")
plt.plot([0, 1], [0, 1], linestyle="--", label="Perfect calibration")
plt.title("Calibration Curve (Test) - SVM Calibrated")
plt.xlabel("Mean predicted confidence")
plt.ylabel("Fraction correct")
plt.grid(True, alpha=0.3)
plt.legend()

out_path = os.path.join(FIG_DIR, "svm_calibrated_calibration_curve.png")
plt.savefig(out_path, dpi=200, bbox_inches="tight")
plt.show()
print("Saved:", out_path)
```



Saved: /content/drive/MyDrive/SkinAI_Project/reports/figures/svm_calibrated_calibration_curve.png


```
from sklearn.model_selection import learning_curve
import numpy as np
import matplotlib.pyplot as plt

train_sizes, train_scores, cv_scores = learning_curve(
    estimator=LinearSVC(class_weight="balanced", random_state=42),
    X=X_train_tfidf,
    y=y_train,
    train_sizes=np.linspace(0.2, 1.0, 6),
    cv=3,
    scoring="accuracy"
)

train_mean = train_scores.mean(axis=1)
cv_mean = cv_scores.mean(axis=1)

plt.figure(figsize=(7,5))
plt.plot(train_sizes, train_mean, marker="o", label="Train Accuracy")
plt.plot(train_sizes, cv_mean, marker="o", label="CV Accuracy")
plt.title("Learning Curve - LinearSVC (Accuracy)")
plt.xlabel("Training samples")
plt.ylabel("Accuracy")
plt.grid(True, alpha=0.3)
plt.legend()

out_path = os.path.join(FIG_DIR, "svm_learning_curve.png")
plt.savefig(out_path, dpi=200, bbox_inches="tight")
plt.show()
print("Saved:", out_path)
```



Saved: /content/drive/MyDrive/SkinAI_Project/reports/figures/svm_learning_curve.png

```
import joblib
import os

os.makedirs(MODELS_DIR, exist_ok=True)

MODEL_PATH = os.path.join(MODELS_DIR, "skin_text_classifier_tfidf_svm_calibrated.joblib")
VECT_PATH = os.path.join(MODELS_DIR, "vectorizer_tfidf.joblib")

joblib.dump(svm_cal, MODEL_PATH)
joblib.dump(vectorizer, VECT_PATH)

print("Saved classifier to:", MODEL_PATH)
print("Saved vectorizer to :", VECT_PATH)
```

Saved classifier to: /content/drive/MyDrive/SkinAI_Project/models/skin_text_classifier_tfidf_svm_calibrated.joblib
 Saved vectorizer to : /content/drive/MyDrive/SkinAI_Project/models/vectorizer_tfidf.joblib

```
def predict_with_confidence(model, vectorizer, text: str):
    """
    Predict label + confidence for Calibrated SVM.
    confidence = predicted class probability (0..1) if predict_proba exists.
    """
    X = vectorizer.transform([text])
    pred = model.predict(X)[0]

    if hasattr(model, "predict_proba"):
        proba = model.predict_proba(X)[0]
        classes = list(model.classes_)
        conf = float(proba[classes.index(pred)])
        return pred, conf

    return pred, None

samples = [
    "I have pimples and red bumps with white pus on my cheeks and forehead.",
    "My skin is very dry and itchy with red patches that crack sometimes.",
    "I have a circular itchy rash with raised red border and clear center."
]

for s in samples:
    p, c = predict_with_confidence(svm_cal, vectorizer, s) # ✅ use svm_cal here
    print("\nTEXT:", s)
    print("PRED:", p)
    print("CONF:", "N/A" if c is None else f"{c*100:.2f}%")
```

TEXT: I have pimples and red bumps with white pus on my cheeks and forehead.

```
PRED: Folliculitis  
CONF: 56.33%
```

```
TEXT: My skin is very dry and itchy with red patches that crack sometimes.  
PRED: Folliculitis  
CONF: 12.35%
```

```
TEXT: I have a circular itchy rash with raised red border and clear center.  
PRED: Folliculitis  
CONF: 15.18%
```

```
import numpy as np  
  
def predict_topk(model, vectorizer, text: str, k: int = 3):  
    X = vectorizer.transform([text])
```